

Approximated PCA

Rodrigo Arias Mallo

July 30, 2017

4 1 PCA

5 1.1 Introduction

6 The PCA method transforms a set of observations of possibly correlated
7 variables into a set of values of linearly uncorrelated variables called principal
8 components using an orthogonal transformation (rotations and reflections).

The transformation, projects the data in a new subspace, in which each new variable it's now uncorrelated. That means that the covariance of each pair of new variables is zero. The main steps of PCA can be summarized as follows:

- Take a dataset X of n variables.
 - Scale and center the variables.
 - From X compute the $n \times n$ covariance matrix S .
 - **Compute the eigenvalues and eigenvectors of S .**
 - *Optional: Ignore some eigenvectors.*
 - Generate a new basis from the selected eigenvectors.
 - Project X into the new basis.

20 Computing the eigenvectors is the principal step of PCA, and different methods can be used.
21

22 1.2 Algorithms

To compute the transformation, different approaches can be taken. In a first attempt, the covariance matrix will be used. Let x_{ij} be the observation j of the variable i . Let n be the number of variables and m the number of observations. Each element s_{ij} in the covariance matrix S is computed by

$$s_{ij} = \frac{\sum x_{ik}x_{jk} - \sum x_{ij}x_{jk}}{n(n-1)}$$

23 Once the covariance matrix S is obtained, the eigenvalues and eigenvectors
 24 need to be computed. One fast and estable method consists in transform the
 25 covariance matrix in a tridiagonal matrix, and then in a diagonal matrix, as
 26 follows:

- 27 1. Take the $n \times n$ target matrix $A = S$.
- 28 2. Compute a tridiagonal matrix T : $A = PTP^T$.
- 29 3. From T compute a diagonal matrix D : $T = QDQ^T$.
- 30 4. The eigenvalues of A are in the diagonal of D .
- 31 5. Compute the eigenvectors from D .

32 Computing a tridiagonal matrix is called **tridiagonalization**. For the di-
 33 agonal, **diagonalization**. Several algorithms exists for both steps.

34 1.2.1 Tridiagonalization algorithms

These algorithms transform a **symmetric** matrix A into a new pair of ma-
 trixes P and T such that P is orthogonal, T is tridiagonal, and $A = PTP^T$

$$\begin{pmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{pmatrix} = P \begin{pmatrix} t_{11} & t_{12} & & \\ t_{21} & t_{22} & t_{23} & \\ & t_{32} & t_{33} & t_{34} \\ & & t_{43} & t_{44} \end{pmatrix} P^T$$

35 Some algorithms can be used for this factorization:

	Algorithm	Complexity	Iterative	Stability
36	Householder	$O(4n^3/3)$	No	Great
	Givens	$O(kn^3)$	No	Good
	Lanczos	$O(kpn^2)$	Yes	Bad
	Others			

37 Where $n \times n$ is the dimension of the matrix A , k is some constant, and p
 38 the number of iterations.

39 1.2.2 Diagonalization algorithms

These algorithms take a **tridiagonal** matrix T into a new pair of matrices
 Q and D such that Q is orthogonal, D is diagonal, and $T = QDQ^T$

$$\begin{pmatrix} t_{11} & t_{12} & & \\ t_{21} & t_{22} & t_{23} & \\ & t_{32} & t_{33} & t_{34} \\ & & t_{43} & t_{44} \end{pmatrix} = Q \begin{pmatrix} d_{11} & & & \\ & d_{22} & & \\ & & d_{33} & \\ & & & d_{44} \end{pmatrix} Q^T$$

⁴⁰ The matrix D contains the **eigenvalues** in the diagonal. Some algorithms
⁴¹ can be used to compute the diagonalization:

Algorithm	Complexity	Convergence
QR	$O(6n^3)$	Cubic
Divide and conquer	$O(8n^3/3)$	Quadratic
⁴² Jacobi	$O(n^3)$	Quadratic
Power iteration	$O(n^3)$	Linear
Inverse iteration	$O(n^3)$	Linear
Others		

⁴³ All these algorithms are iterative and $n \times n$ is the dimension of the matrix A .
⁴⁴ The proposed algorithms for the task are the Householder followed by a
⁴⁵ QL algorithm, which is a variation of the QR algorithm.

⁴⁶ 1.3 Householder

⁴⁷ The Householder tridiagonalization it's a process where a matrix A is trans-
⁴⁸ formed by multiplying with an orthogonal matrix $P^{(k)}$: $P^{(k)} = I - 2ww^T$
⁴⁹ Such matrix $P^{(k)}$ has been prepared, so that $P^{(k)}A$ is a new matrix, with
⁵⁰ zeros below the $k + 1$ element in the k column. This new matrix, has the
⁵¹ same eigenvalues as the previous A . The step is repeated until the final
⁵² matrix has only elements in the diagonal, and the two sub-diagonals. The
⁵³ process is similar to a Gaussian elimination.

⁵⁴ 1.4 QL

⁵⁵ The QL algorithm completes the process of diagonalization of the matrix,
⁵⁶ by iteratively approaching to the solution. There are $n - 1$ stages in a matrix
⁵⁷ of $n \times n$. At each stage i , the elements of the offdiagonals o_i are iteratively
⁵⁸ decreased, until they reach 0. At this moment, the algorithm steps to the
⁵⁹ next stage. A step is defined as one update of o_i and d_i in each iteration.

⁶⁰ 2 Experiments with Householder algorithm

⁶¹ The reduction of bits in the mantisa of the floating points used by the
⁶² algorithm can lead to an acceleration in the ALU. However, the precision of
⁶³ the results can be affected by the number of bits used.

⁶⁴ For this reason, a set of experiments are performed, to test how the error
⁶⁵ grows as the number of bits of the mantisa is reduced.

⁶⁶ The library MPFR is designed to perform computations with an arbi-
⁶⁷ trary mantissa length. After rewrite the Householder algorithm, the results
⁶⁸ can be compared with a golden execution. This golden execution is done

69 with a very big mantissa, such that the error of the result is so low that can
70 be ignored, compared with the errors produced in the experiments.

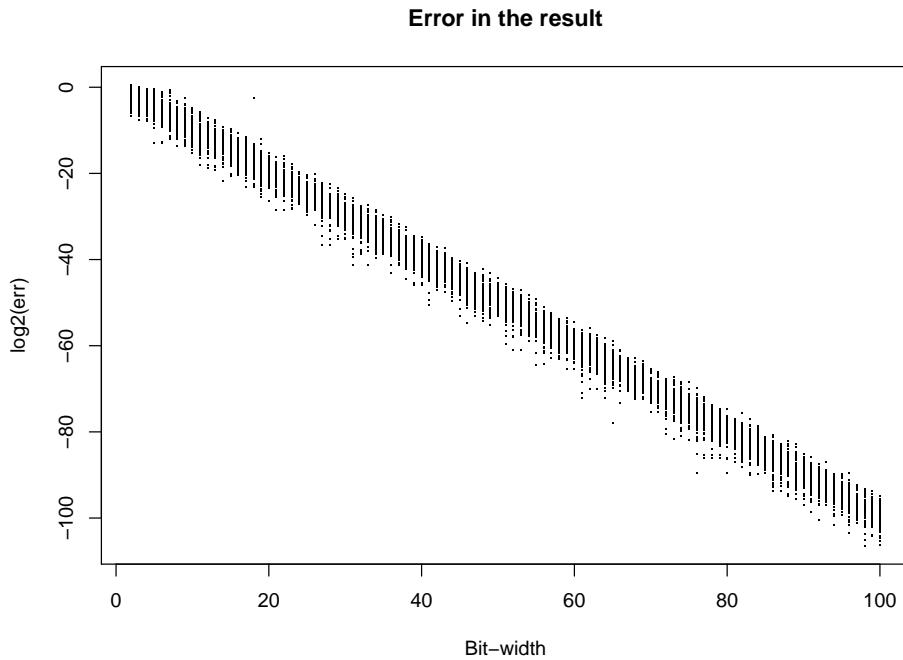
71 The aim of this set of experiments is to reduce the storage size needed
72 while keeping a reasonable error.

73 **Experiment A**

74 **Objective** Understand how the bit-width b affects the error in the House-
75 holder algorithm.

76 **Description** For an input matrix A symmetric and random, of size $n \times n$
77 with $n = 5$, a gold result is computed, using a high precision computation
78 with $b = 500$. Then, in each run, the bit-width b is set to a value in the
79 range $[2, 100]$ and Householder runs again on the same input. The error Δ
80 is measured compared as with the golden result.

81 **Result** The error is drawn as the bit-width b grows.



82 **Conclusion** It can be observed that $\log_2 \Delta$ is close to $-b$. More experi-
83 ments are needed to test this hypothesis.

⁸⁴ **Experiment B**

⁸⁵ **Objective** Test if the relation between the $\log_2 \Delta$ and b is linear.

⁸⁶ **Description** Let $X_b = \log_2 \Delta$ and $Y = b + \bar{X}_b$. We can now plot Y as b grows. If $\bar{X}_b = -b$, then Y should be 0. For each bit-width b in the range ⁸⁷ $[2, 100]$, a set of 10000 simulations are performed with $n = 5$, and the sample ⁸⁸ mean \bar{X}_b is computed.

⁹⁰ **Result** The random variable Y seems to be constant as the bit-width grows, as can be seen in the figure 1.

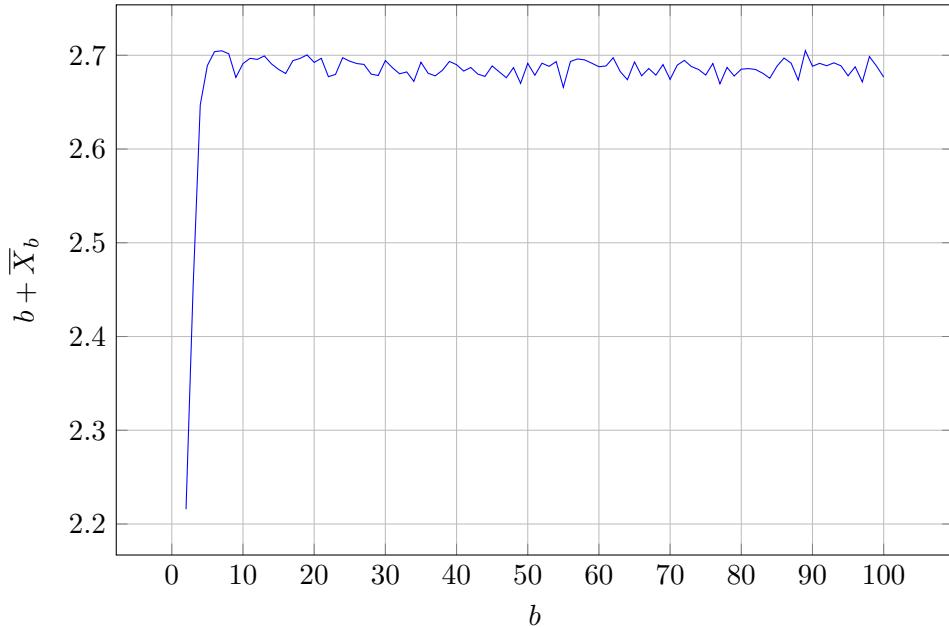


Figure 1: Plot of $b + \bar{X}_b$ as the number of bits b increases.

⁹¹

⁹² **Conclusion** Further statistical analysis reveals that, if the values with ⁹³ $b < 5$ are not considered, Y is independent of b , and has a mean in the ⁹⁴ interval 2.68691 ± 0.00190 at a confidence level of 0.95. Then, the relation ⁹⁵ between \bar{X}_b and b is linear. More experiments are needed to determine if ⁹⁶ the mean of Y is always constant.

97 **Experiment C**

98 **Objective** Test if there is any relation between Y and the precision of
 99 each variable of the Householder algorithm.

100 **Description** The precision is now determined by a vector \mathbf{b} , so that each
 101 variable used in the Householder algorithm is assigned a own b_i . Let \mathbf{v} be the
 102 set of variables. Then v_1 is the matrix A, v_2 the diagonal, v_3 the offdiagonal,
 103 and the others v_i are the internal parameters of Householder algorithm. The
 104 variable v_i is set a bit-width of b_i . The experiment is repeated as in the
 105 experiment B, but we only change one b_i at the time, selecting a precision
 106 from [5,100], while the others are set at 500 bits.

107 **Result** The experiments show that the variable v_4 don't produce any error
 108 in the output. This variable is a internal scale variable of the Householder
 109 algorithm, and will be ignored, to show the other variables.

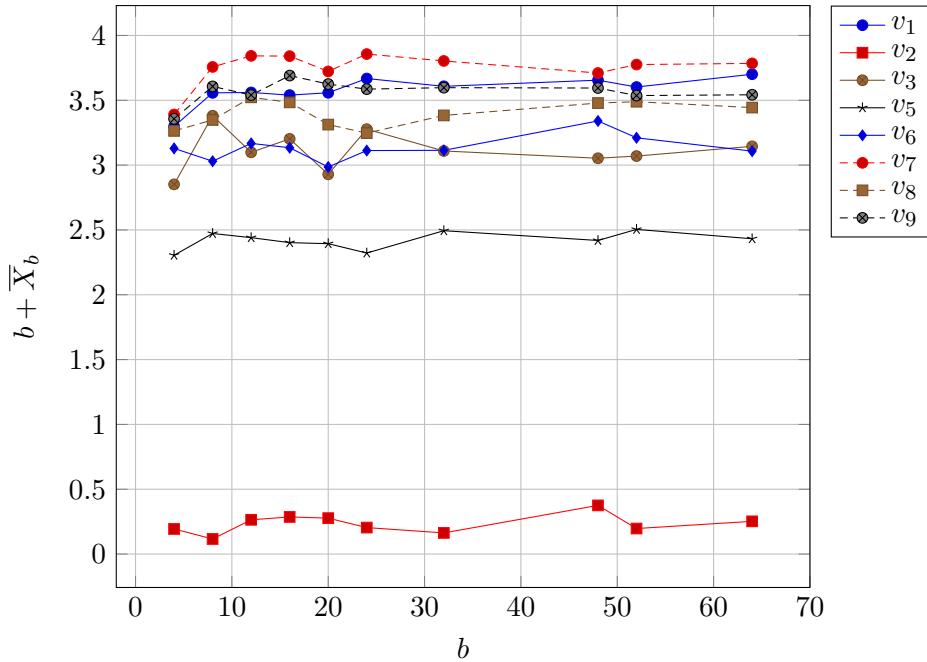


Figure 2: Different error mean for each variable.

110 **Conclusion** It can be shown that the precision of the variables affect the
 111 value of the Y , and also that is independent of the value of b .

112 **Experiment D**

113 **Objective** Test if there is any relation between Y and the condition number
114 κ of the input matrix.

115 **Description** The Householder algorithm is executed with a fixed bit-
116 width $b = 50$ and compared with the gold result with $b = 500$. The error is
117 plotted against the logarithmic condition number $\log \kappa$ computed from the
118 input matrix.

119 **Result** In the figure 3 it can be seen that there is no clear relationship
between the condition number and Y .

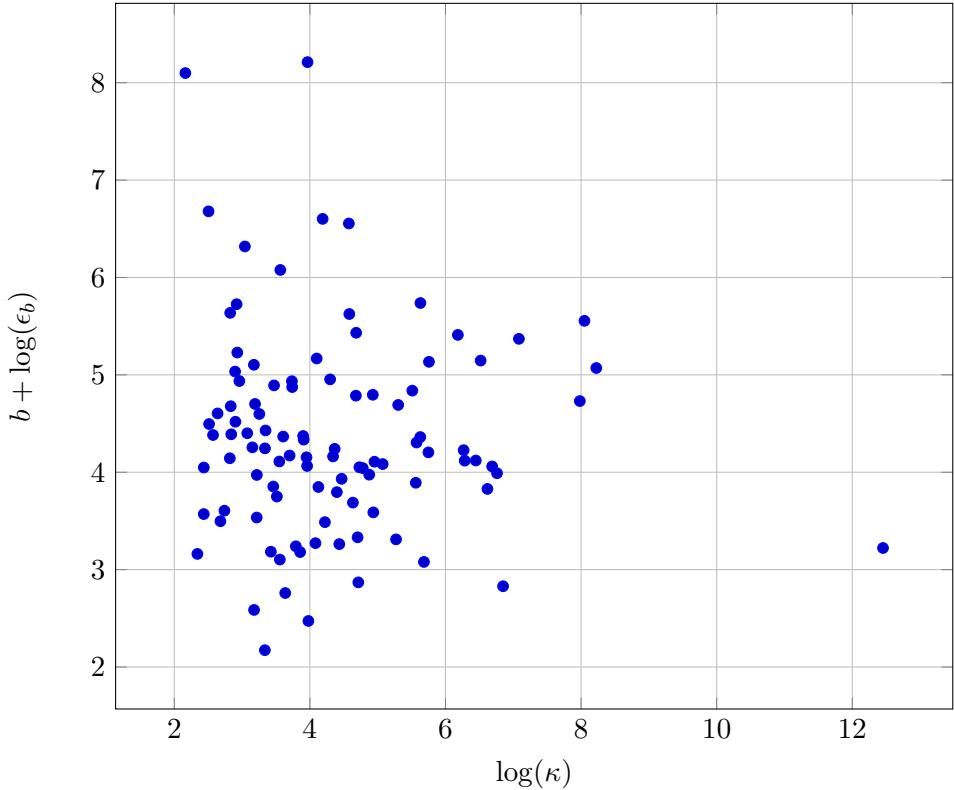


Figure 3: Using a fixed bit-width $b = 50$ the error is plotted against the condition number.

120

121 **Conclusion** The condition number doesn't affect the random variable Y
122 in a meaningful way.

123 **Experiment E**

124 **Objective** Determine the relation between Y and n .

125 **Description** The Householder algorithm is executed varying both the bit-width b and the size of the $n \times n$ input matrix. The error is measured by 126 computing the norm-2 of the difference between the output vectors of the 127 gold result, and the current result.

129 **Result** The mean of the sample Y is computed from the values with the 130 same input size n . In the figure 4, the cubic mean \bar{Y}^3 is plotted against n .

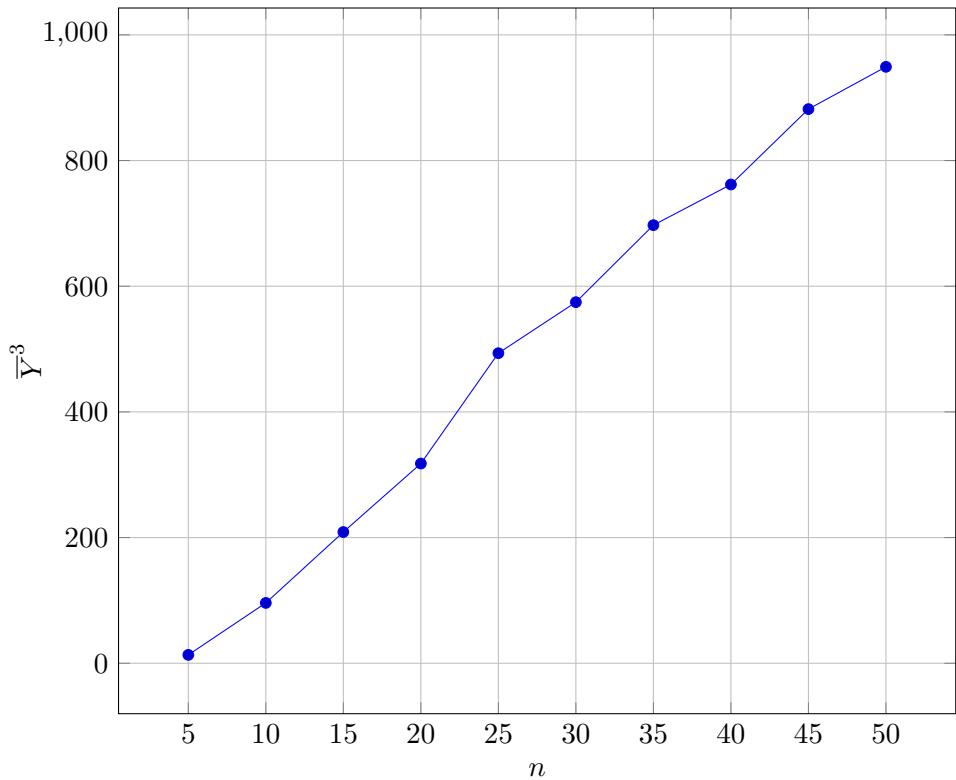


Figure 4: The error \bar{Y}^3 is computed with varying matrix sizes n .

131 **Conclusion** It can be seen that the relation between \bar{Y}^3 and n seems 132 linear, but more data is necessary. A new experiment should be designed to 133 cope with bigger matrices.

¹³⁴ **Experiment F**

¹³⁵ **Objective** Determine the relation between Y and n when n is big.

¹³⁶ **Description** The experiment E is now designed to deal with bigger matrices, and is executed with values of n in the range of [500, 5000]. To compute the exact exponent, we plot the ratio $Y/\log_2 n$ as n grows.

¹³⁹ **Result** It can be seen that the ratio is almost constant, as n grows, and the mean is 2.78857, represented as a red line in the figure 5.

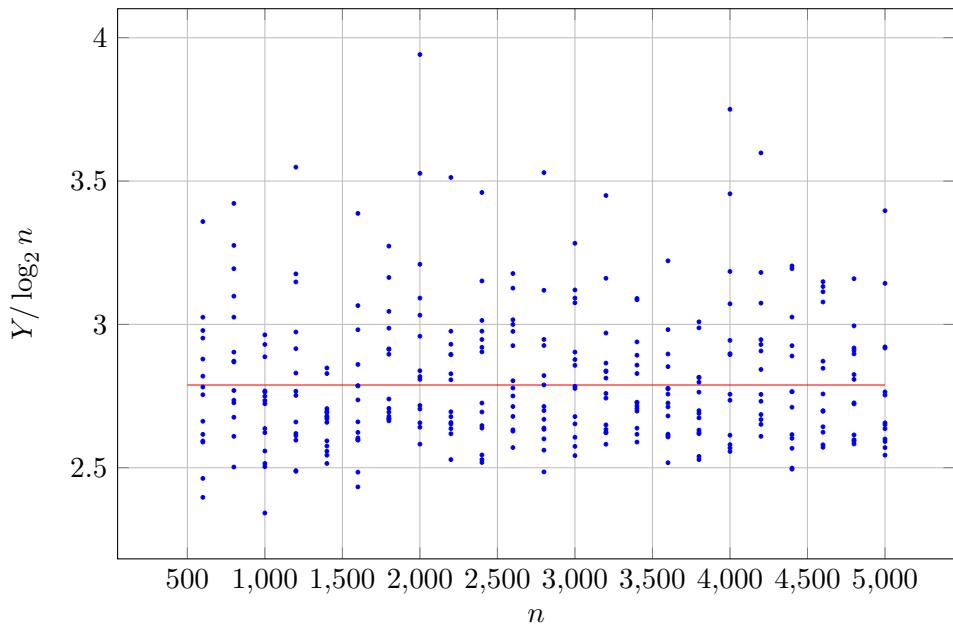


Figure 5: The error \bar{Y} with varying matrix sizes n .

¹⁴⁰

Conclusion The error Y seems to grow with n , with a power of $\alpha = 2.78857$,

$$Y/\log_2(n) \approx 2.78857 = \alpha$$

So, we can compute an approximation of the rounding error Δ as a function of b and n , for bigger values of n .

$$X = \log_2(\Delta) \approx -b + Y = -b + \alpha \log_2(n)$$

And as $\epsilon = 2^{-b}$, we get:

$$\Delta \approx \epsilon \cdot n^\alpha$$

141 **Experiment G**

142 **Objective** Determine if we can use less storage space while maintaining a
143 low error, with different precisions in the variables of Householder algorithm,
144 as well as the input matrix A , and the two vectors diagonal and offdiagonal.

145 **Description** In this experiment the variables used in the Householder al-
146 gorithm are set with individual precisions. The input matrix A and the
147 diagonal vectors diagonal and offdiagonal are set each with individual pre-
148 cisions. The golden result is computed using 500 bits of precision in all the
149 variables. The storage size is plotted against the error.

150 **Result** It can be seen in the figure 6 the relation between the error and
151 the storage size in bits. When the main variables like the input matrix A , or
152 the diagonal and offdiagonal vectors are assigned a low precision, less space
153 is needed, but the error increases. Also, by using a mix of high and low
154 precision variables, the error keeps high.

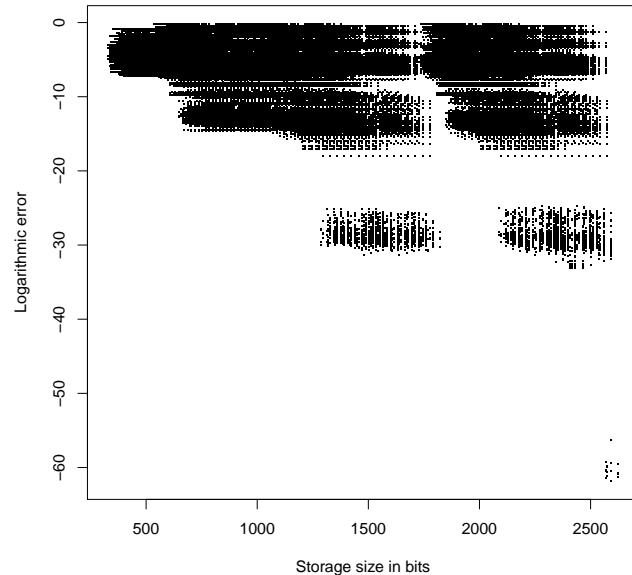


Figure 6: The error \bar{Y} with varying matrix sizes n .

155 **Conclusion** It seems better to select all precision to the same value, in
156 order to obtain the smallest error.

157 **Experiment H**

158 **Objective** Determine the effect of the precision in the different elements in
159 the input matrix A and the vector diagonal and offdiagonal, as they change.

160 **Description** In each step, a random configuration of bits is assigned to
161 each element in the matrix and in both vectors. The possible values are
162 $\{8, 16, 32, 64\}$, and are selected randomly with a probability of $1/4$ each.

163 Once the configuration of bits for the data storage is completed, the
164 Householder algorithm runs with the internal variables at 64 bits. Then,
165 the error is computed from the gold result.

166 **Result** It can be seen, in red, the configurations now have a big error and
167 use more space than the previous experiment.

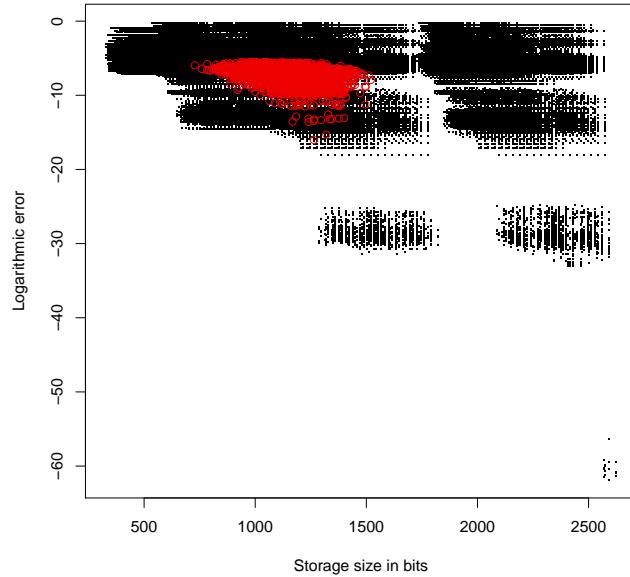


Figure 7: The error \bar{Y} with varying matrix sizes n .

168 **Conclusion** Is better to left all the precisions inside the matrix or vectors
169 to the same value.

170 3 Experiments with QL

171 Some experiments were designed to reduce the number of steps required.
172 The input matrices used in the experiments were generated using random
173 elements uniformly distributed in $[-1, 1]$. The experiments A, B, C and
174 D were executed using single precision (float), while the experiment E is
175 executed using variable precision with the MPFR library.

The errors are measured by comparing the diagonal d' after the QL algorithm (which should contain the eigenvalues), with the diagonal d of the modified QL. The values of diagonals are first sorted, and the relative error δ is then measured as:

$$\delta = \frac{\|d - d'\|_2}{\|d'\|_2}$$

176 Experiment A

177 **Objective** Determine if we can reduce the number of steps in QL, while
178 keeping the relative error low.

Description First an estimation is used to compute the relative error at each step of the QL algortihm. Then the algorithm is stopped when the estimation measures a error below a threshold. The estimation is computed over the offdiagonal elements $o_i = a_{i,i+1}$ and the diagonal elements $d_i = a_{ii}$ as

$$s = \frac{\sum_{i=1}^{N-1} |o_i|}{\sum_{i=1}^N |d_i|}$$

179 **Result** In the figure 8 there can be seen 3 executions with random matrices
180 of size $n \times n$ with $n \in \{50, 100, 200\}$. The relative error is compared with the
181 ratio of completion of the QL algorithm. The error decreases more quickly
182 in the last steps, and the slope becomes sharper as the size of the matrix
183 increases. The 2 lines in red mark a interesting region, where the error is
184 below 1% and the ratio is less than 90%.

185 **Conclusion** There is no enough reduction of the number of steps to consider this technique useful.
186

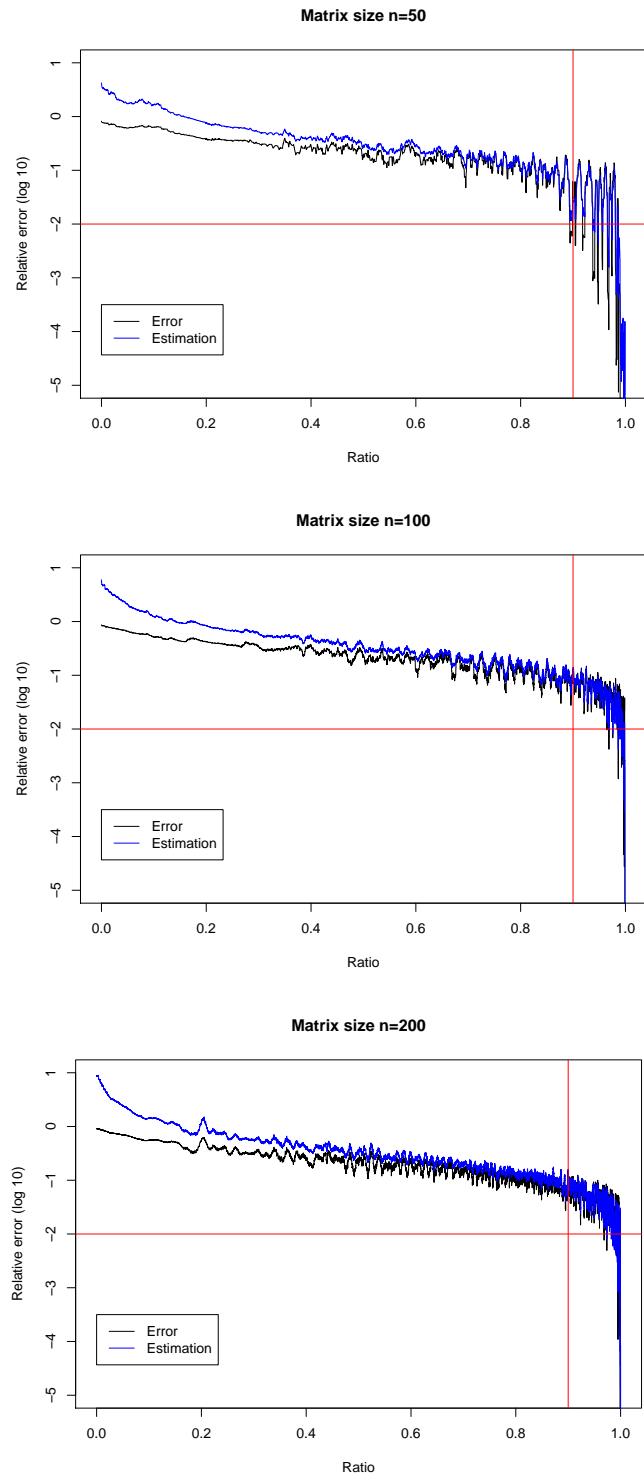


Figure 8: The logarithmic relative error as the number of steps grow.

187 **Experiment B**

188 **Objective** Determine if we can reduce the number of steps in QL, while
189 keeping the relative error low.

190 **Description** The QL algorithm iterates over the elements in the diagonal
191 and subdiagonal, until the first non-zero subdiagonal element becomes zero.
192 Then it moves to the next non-zero element. The number of iterations can
193 vary, but a maximum of $m = 30$ is considered enough for any case. In this
194 experiment, the maximum number of iterations m is set to lower value, and
195 if is reached, the execution moves to the next non-zero element.

196 **Result** In the figure 9 the error is measured against the ratio, with differ-
197 ent values for m . Again a region of interest is represented in red lines, with
198 less than 1% of relative error and less than 90% of steps. The red points
199 have $m = 2$, and they are almost all in the interesting region.

200 More runs are tested in figure 10 with a fixed value of $m = 2$, with input
201 matrix sizes of $n \times n$ with $n \in \{50, 100, 200\}$ an respective number of runs
202 $r \in \{10^5, 10^4, 10^3\}$. The mean ratio seems to grow slowly with n , but stays
203 close to 65%.

204 **Conclusion** By using this technique, the number of steps can be reduced
205 down to 65%, while the relative error is mostly below 1%.

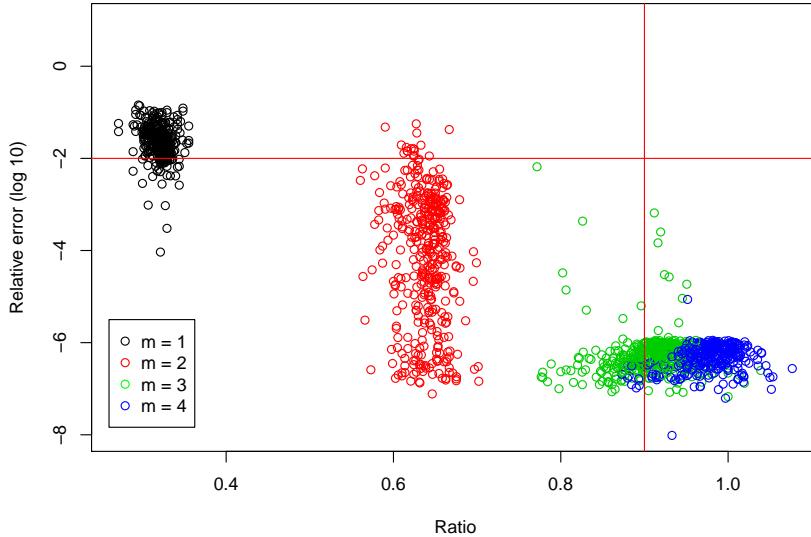


Figure 9: The logarithmic relative error with different values of m .

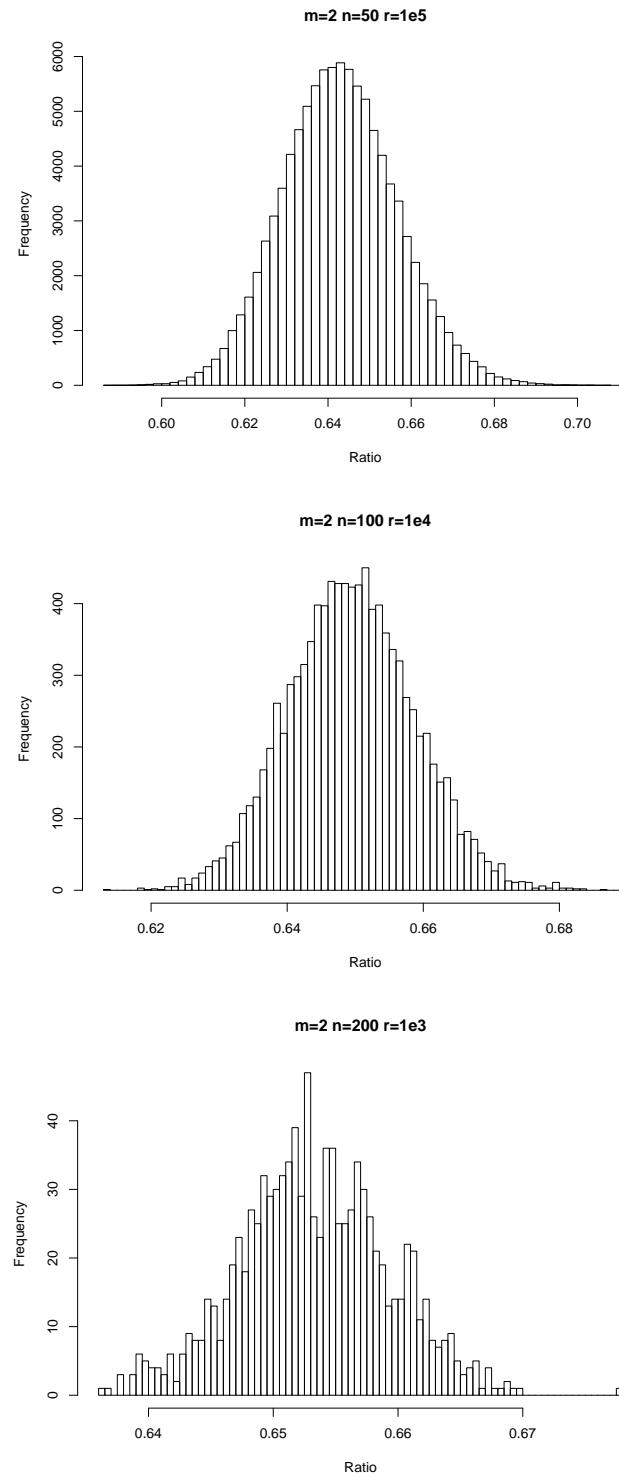


Figure 10: The histogram of the ratio as n grows, with $m = 2$, using r simulations.

206 **Experiment C**

207 **Objective** Determine if we can reduce the number of steps in QL, while
208 keeping the relative error low.

209 **Description** In this experiment, the comparison of the subdiagonal ele-
210 ments with 0 is relaxed by introducing a small tolerance t . So that a subdi-
211 agonal element o_i is considered 0 if $|d_i| + |d_{i+1}| + |o_i| \leq (1+t)(|d_i| + |d_{i+1}|)$.

212 Different values of t are tested while recording the ratio of steps and the
213 relative error.

214 **Result** In the figure 11 different tolerances are tested for random input
215 matrices of size 50×50 , while the ratio and logarithmic relative error are
216 measured. Each color represents a different tolerance, as can be seen in the
217 figure 12. A horizontal red line corresponds to 1% of relative error, and it
218 can be seen that there are some values in the interesting region, below this
219 line.

220 **Conclusion** The experiment shows a interesting result, with a ratio of
221 30% and 50% we have relative errors between 10^{-2} and 10^{-6} . The number
222 of steps can be reduced down to 30% while keeping a relative error close to
223 1%.

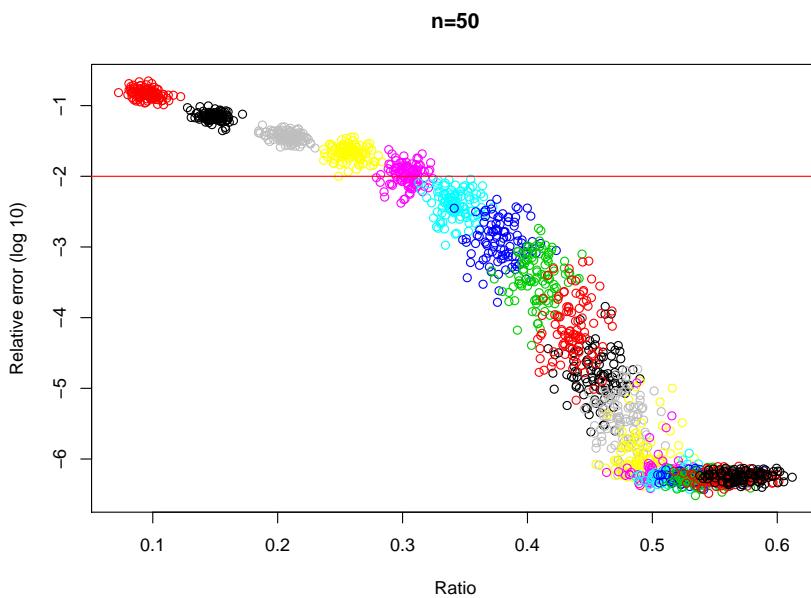


Figure 11: The logarithmic relative error and ratio with different tolerances.

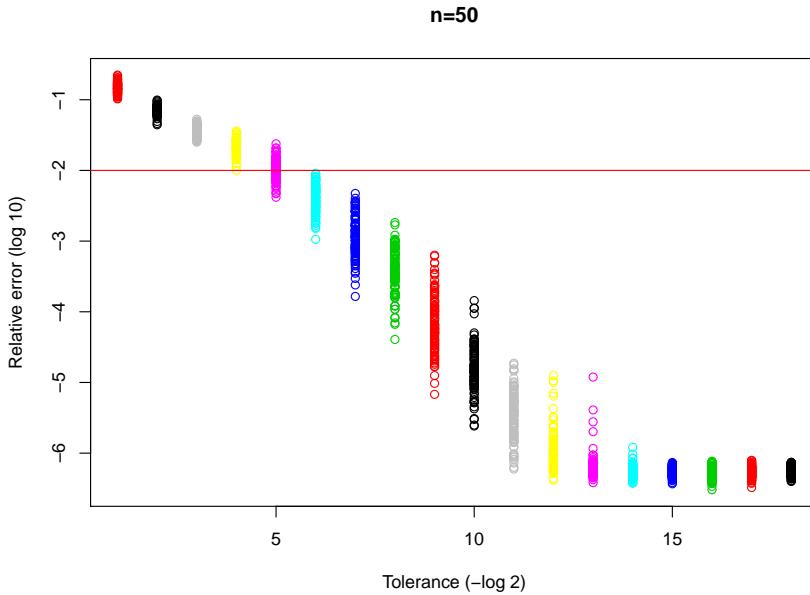


Figure 12: The logarithmic relative error with different tolerances.

224 **Experiment D**

225 **Objective** Determine if we can reduce the number of steps in QL, while
226 keeping the relative error low.

227 **Description** The technique of loop perforation is tested in the following
228 way: we skip each step with probability p . Different values of p are tested
229 with $\log_{10} p \in [-7, -1]$. The size $n \times n$ of the random input matrices is also
230 tested with $n \in \{10, 20, 50, 100\}$. The relative error and the ratio are finally
231 measured.

232 **Result** In the figure 13 the relative error is plotted against the ratio, each
233 color represent a different value of p . Points in the interesting region, below
234 1% of relative error and less than 90% of ratio, are never achieved. In the
235 figure 14 it can be seen which is the effect of the probability p on the error.
236 Same colors are used in both figures.

237 **Conclusion** This technique doesn't seem to provide interesting results,
238 it even introduces more error and more required steps than the original
239 method.

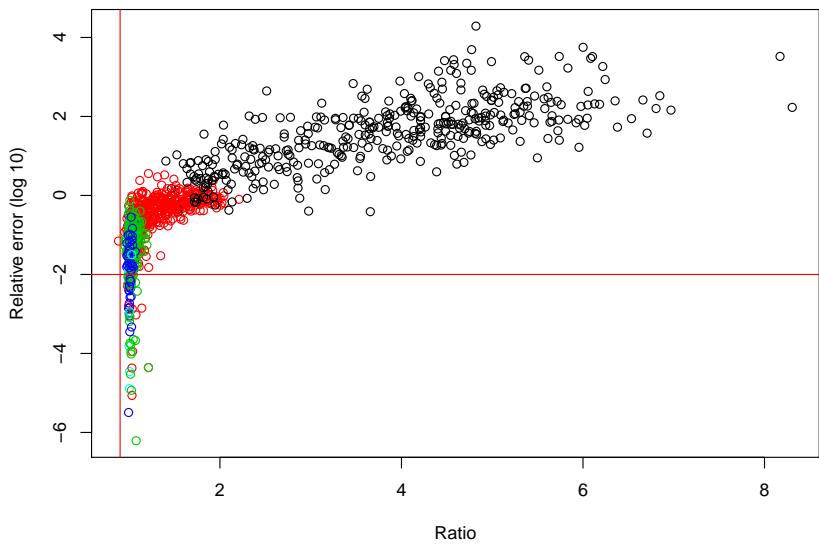


Figure 13: The logarithmic relative error and ratio with different probabilities.

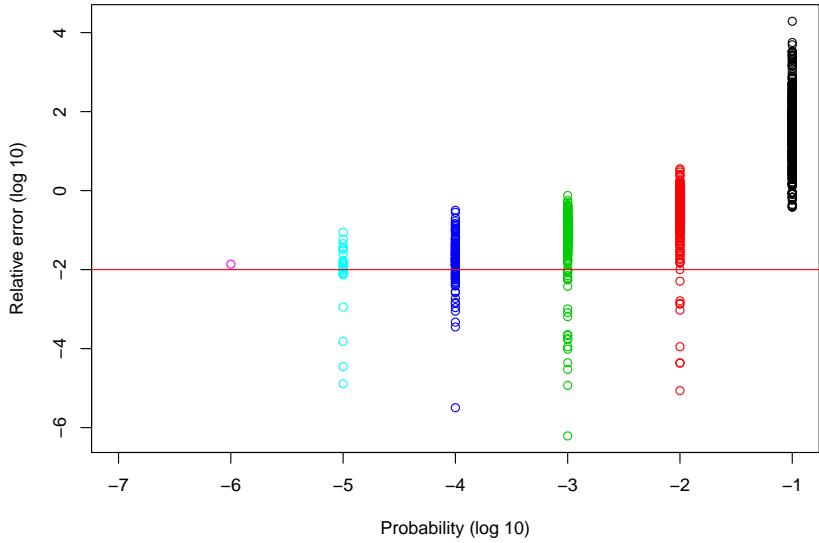


Figure 14: The logarithmic relative error with different probabilities.

240 **Experiment E**

241 **Objective** Determine if the relative error can be reduced while the precision
 242 of the variables in the QL algorithm are changed.

243 **Description** The significant is set to b bits of precision with $b \in \{2, 6, \dots, 98\}$,
 244 then the results are compared with a reference result, with $b = 500$. The
 245 relative error and the ratio are measured. The ratio is computed based on
 246 the current execution steps over the steps with $b = 98$. Different matrix
 247 sizes of $n \times n$ with $n \in \{10, 20, 50, 100\}$ are also tested.

248 **Result** In the figure 15 the error is shown as the number of bits b increases.
 249 The error decreases almost linearly with b , the same happened with
 250 Householder. To see the effect of the different sizes, let $Y = b + \log_2 \delta$, as
 251 shown in the figure 16. It can be seen that as the size n grows, the variable
 252 Y increases. Finally, in the figure 17 the ratio and logarithmic relative error
 253 are plotted, as b grows. In the interesting region, delimited by a 1% relative
 254 error, and a 90% ratio, there are some interesting points. However, as the
 255 size increases, the points move outside the interesting region.

256 **Conclusion** For relative small sized matrices, this technique can reduce
 257 the number of steps up to 60%, while keeping an relative error lower than
 258 1%. However as the matrices grow, the ratio increases, as well as the error.

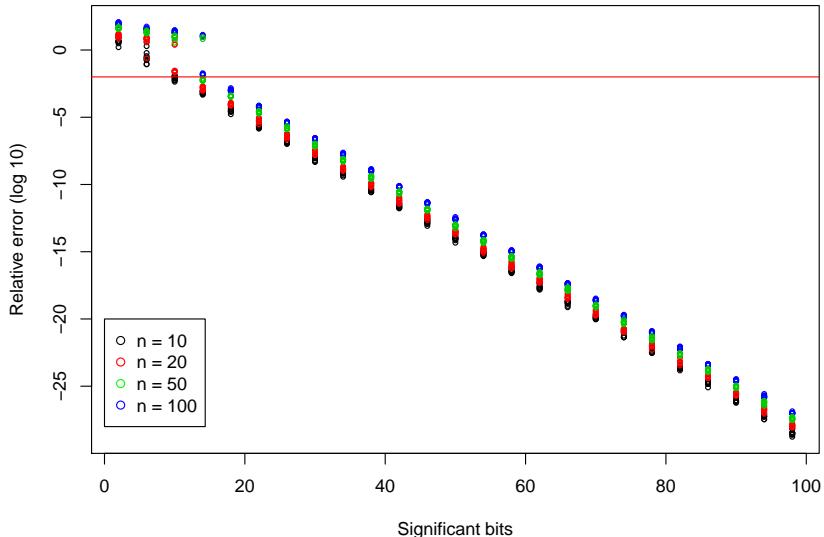


Figure 15: The logarithmic relative error as b grows.

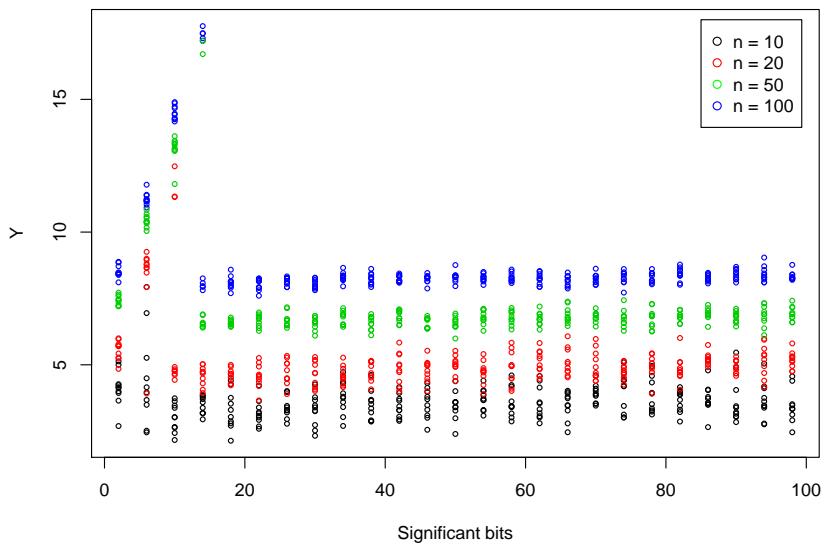


Figure 16: The difference $b + \log_2 \delta$ as b grows.

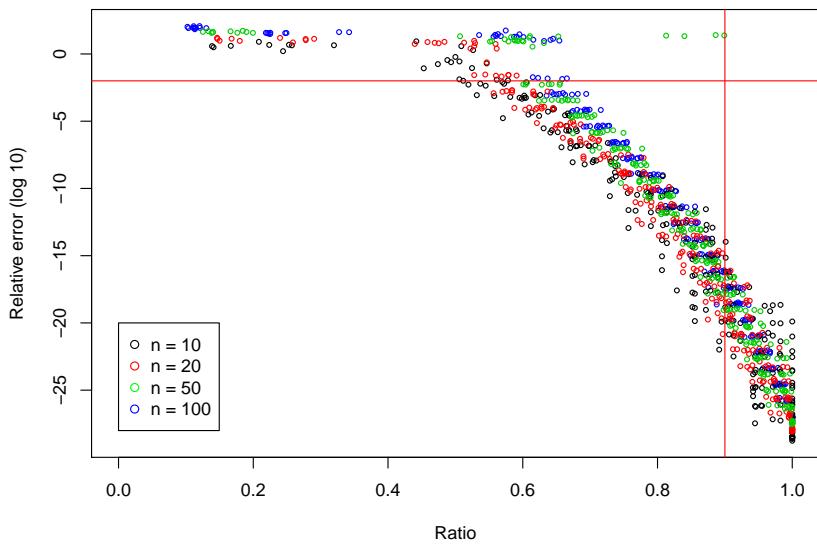


Figure 17: The $\log_{10} \delta$ and ratio as b grows.

259 **4 Discussion**

260 The experiments in the Householder algorithm reveal that the relation be-
261 tween $\log_2 \Delta$ and b is linear, and we can estimate the error by $\Delta \approx \epsilon n^\alpha$.
262 A reduction in the storage size can be achieved with a reduction in b . The
263 best configuration is to keep all the variables at the same precision. The
264 reduction of size is proportional to b and n .

265 In the other hand, by allowing a relative error of 1% in the QL algorithm,
266 different techniques can be used to reduce the number of steps. The best ones
267 are: The introduction of a tolerance in the measurement of the subdiagonal
268 elements, which can reduce the number of steps down to 35%. The reduction
269 of b , to around 50% – 70%. And the limit in the maximum iterations, which
270 at $m = 2$ reduces the steps to 60% – 65%.